


DEEP LEARNING

LECTURE 6: GENERATIVE ADVERSARIAL NETWORKS

goorm

KAIST AI
Graduate School of AI





- Namju Kim. Generative Adversarial Networks (GAN)

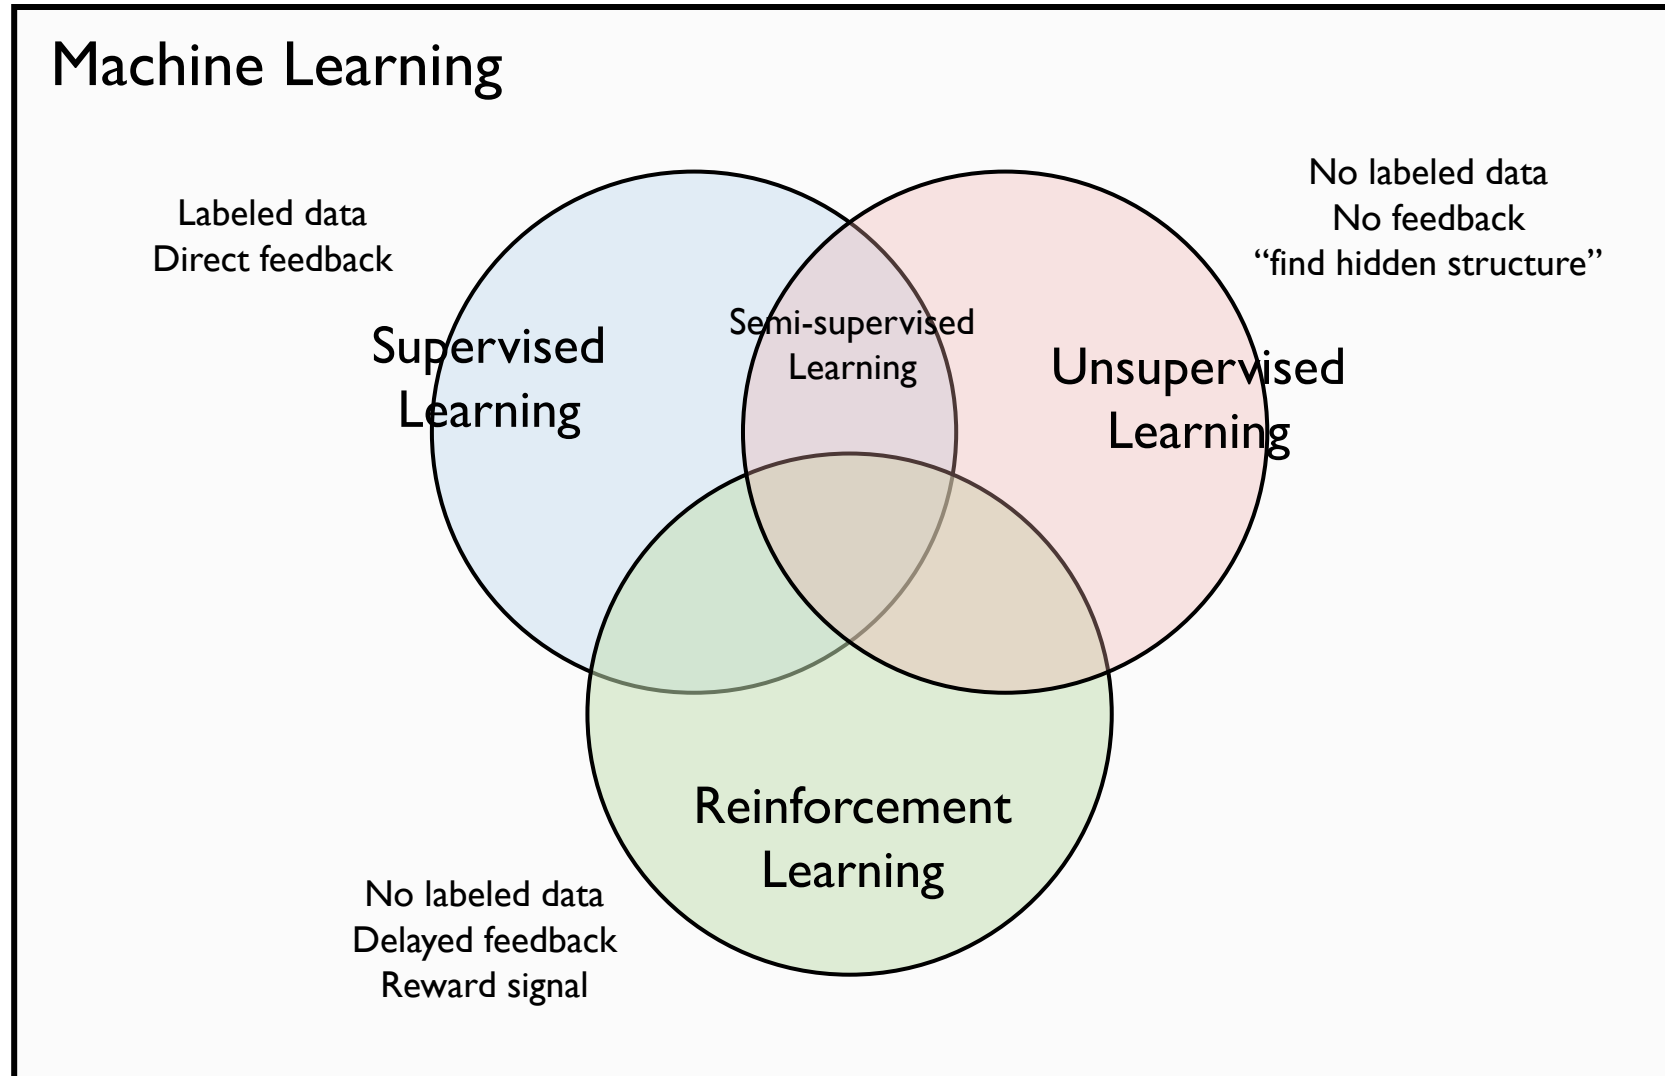
<https://www.slideshare.net/ssuser77ee21/generative-adversarial-networks-70896091>

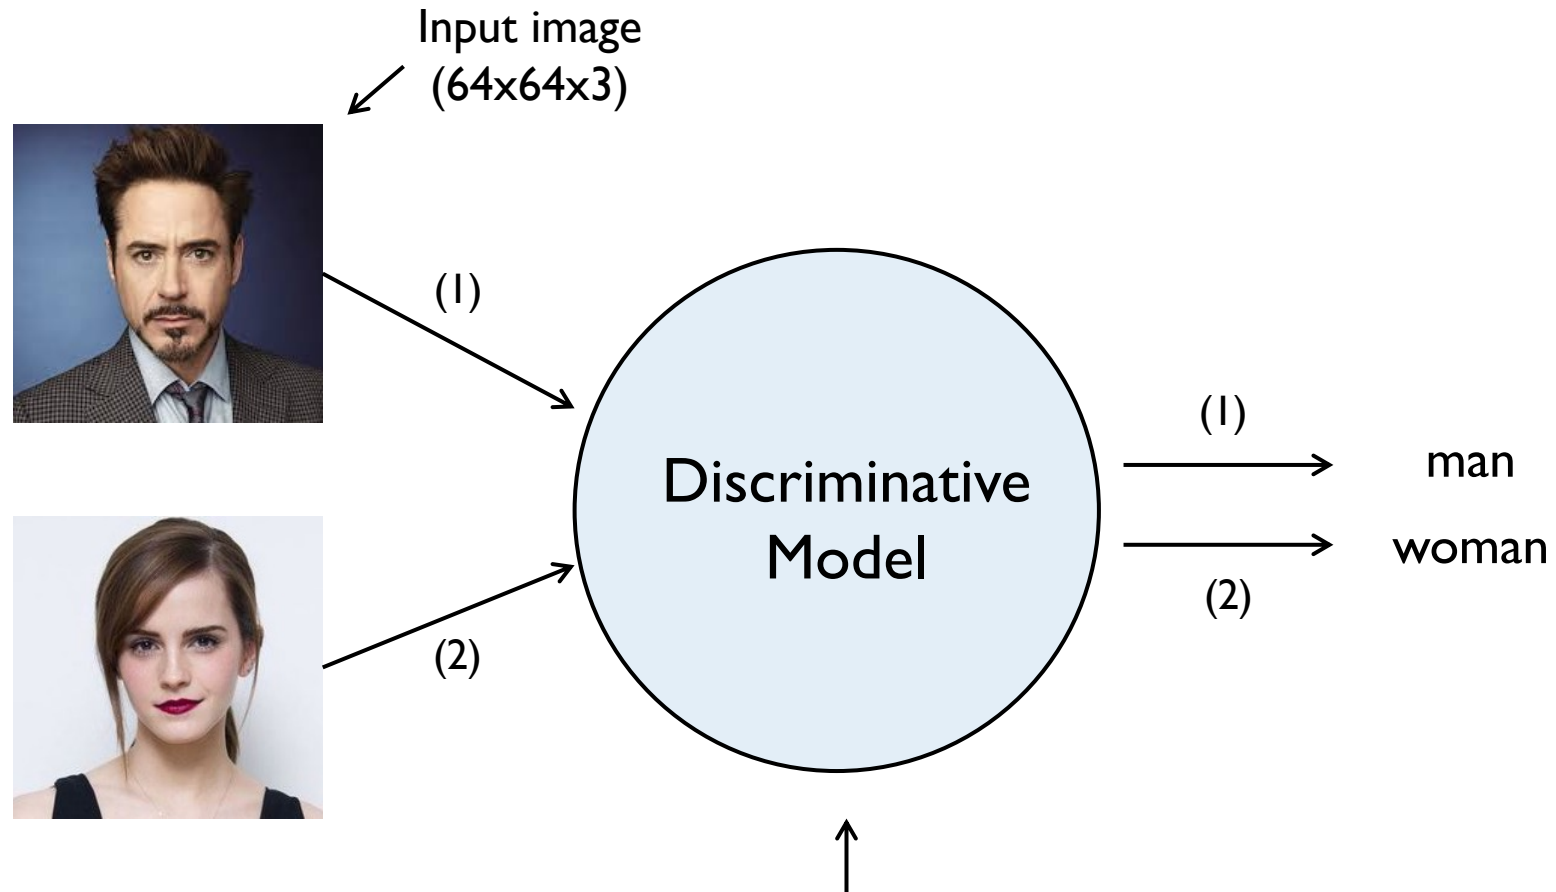
- Taehoon Kim. 지적 대화를 위한 깊고 넓은 딥러닝

<https://www.slideshare.net/carpedm20/ss-63116251>

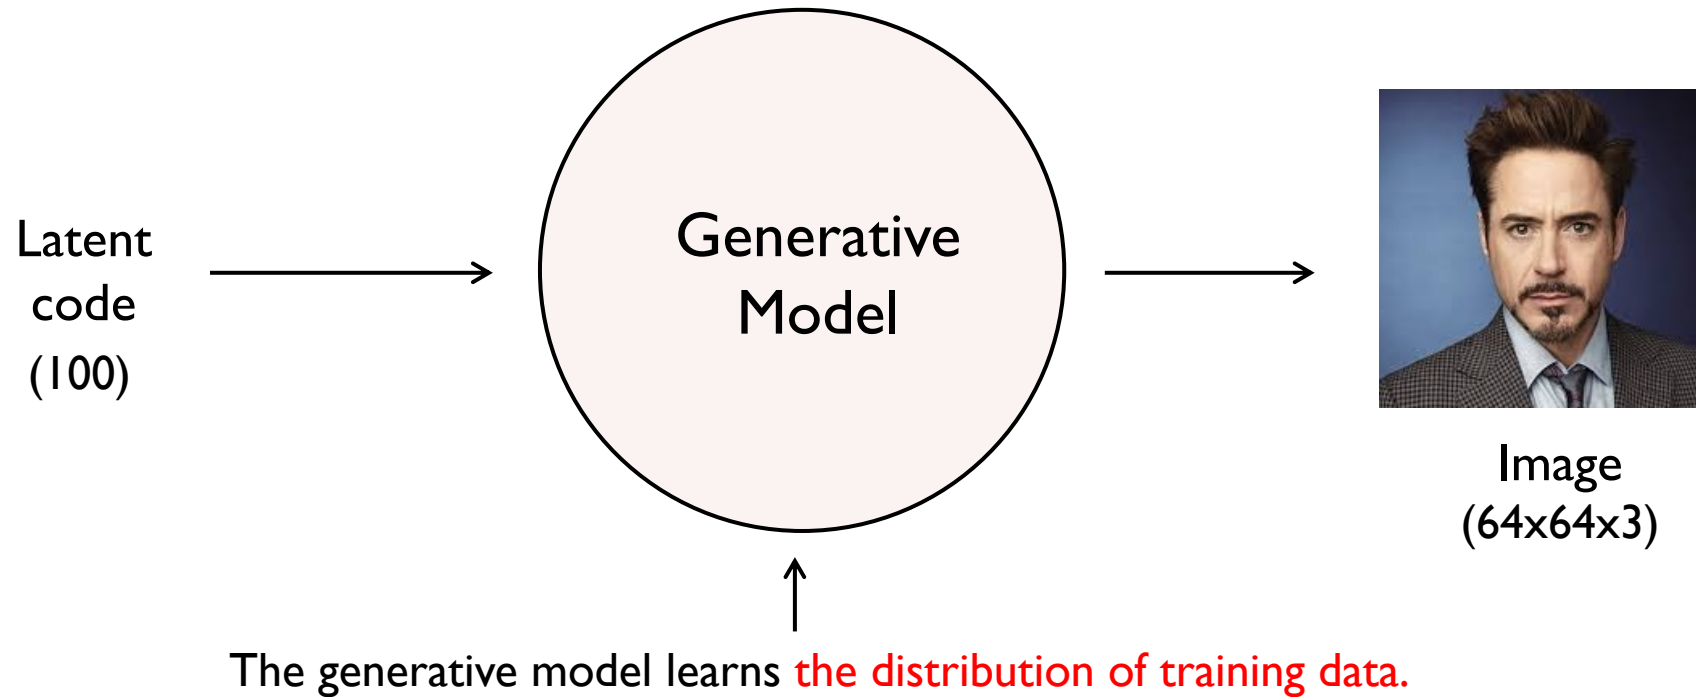
01 Introduction







The discriminative model learns **how to classify** input to its class.



Probability Distribution



Introduction




Probability Basics (Review)



Random variable

X	1	2	3	4	5	6
$P(X)$	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{0}{6}$	$\frac{1}{6}$	$\frac{2}{6}$

Probability mass function





What if x is actual images in the training data?

At this point, x can be represented as a (for example) 64x64x3 dimensional vector.






Probability density function




There is a $p_{data}(x)$ that represents the distribution of actual images.





Let's take an example with human face image dataset.
Our dataset may contain few images of **men with glasses**.




x_1 is a 64x64x3 high dimensional vector
representing **a man with glasses**.





Our dataset may contain many images of **women with black hair**.




x_2 is a 64x64x3 high dimensional vector representing **a woman with black hair**.






Our dataset may contain very many images of **women with blonde hair**.



x_3 is a 64x64x3 high dimensional vector representing **a woman with blonde hair**.



Our dataset may not contain **these strange images**.




x_4 is an 64x64x3 high dimensional vector representing **very strange images**.



The goal of the generative model is to find a $p_{model}(x)$ that approximates $p_{data}(x)$ well.

↗ Distribution of images generated by the model

↘ Distribution of actual images



02

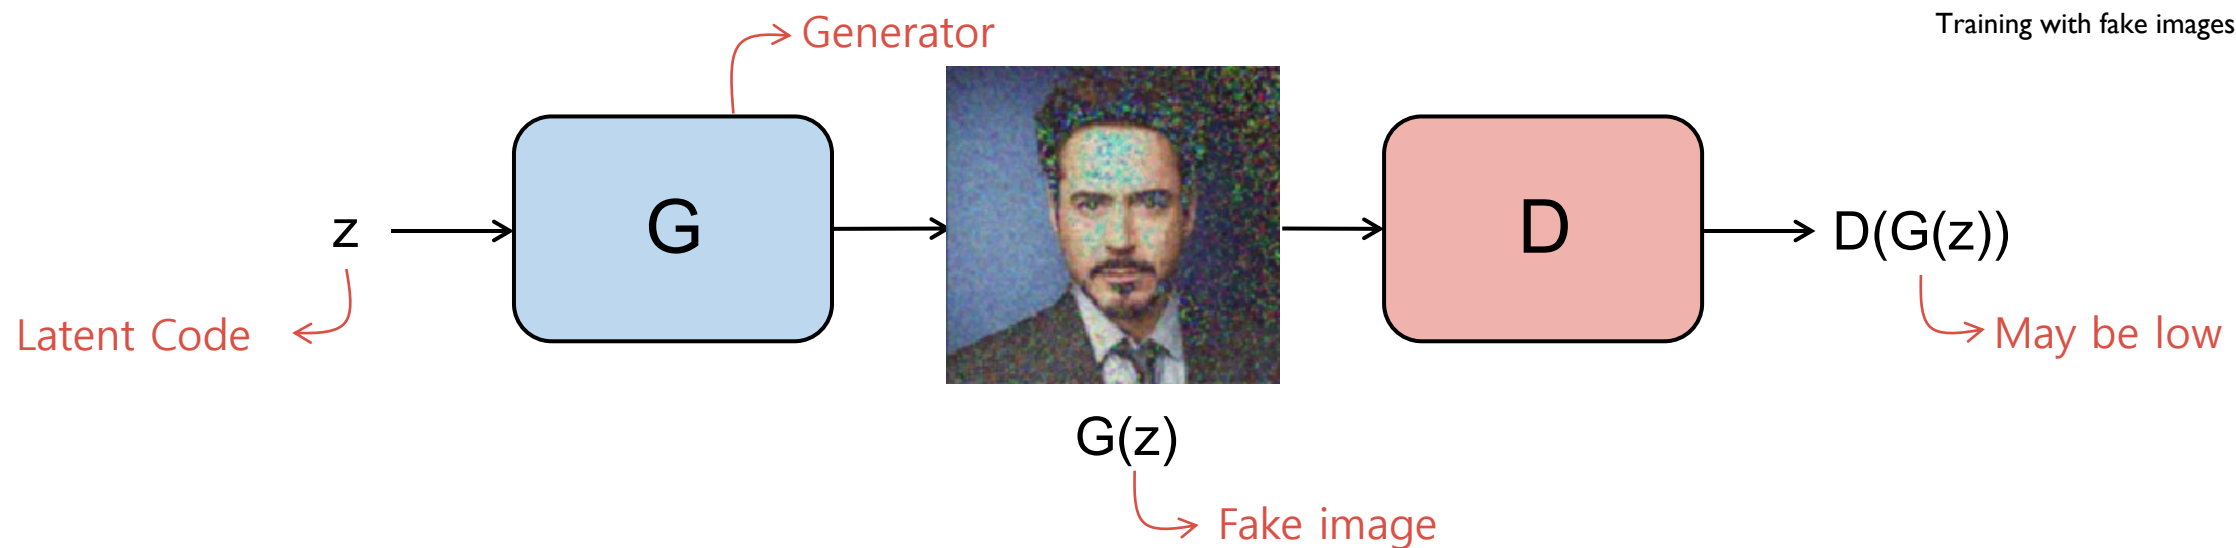
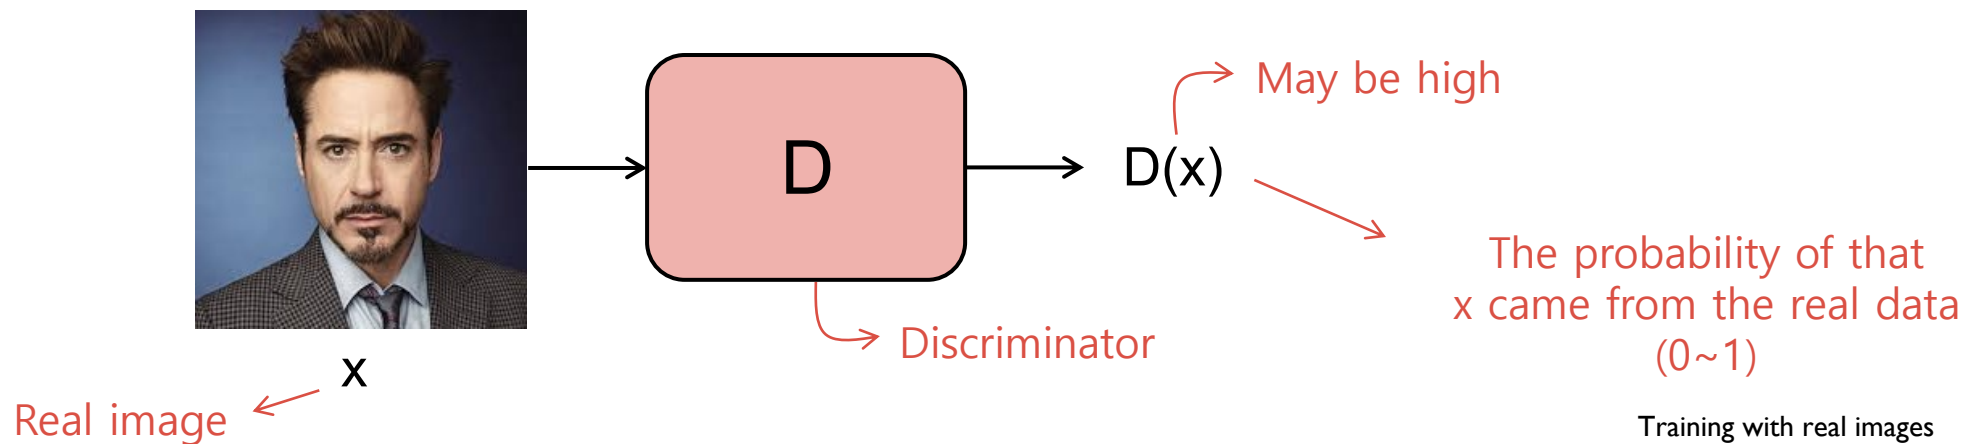
Generative Adversarial Networks



Intuition in GAN



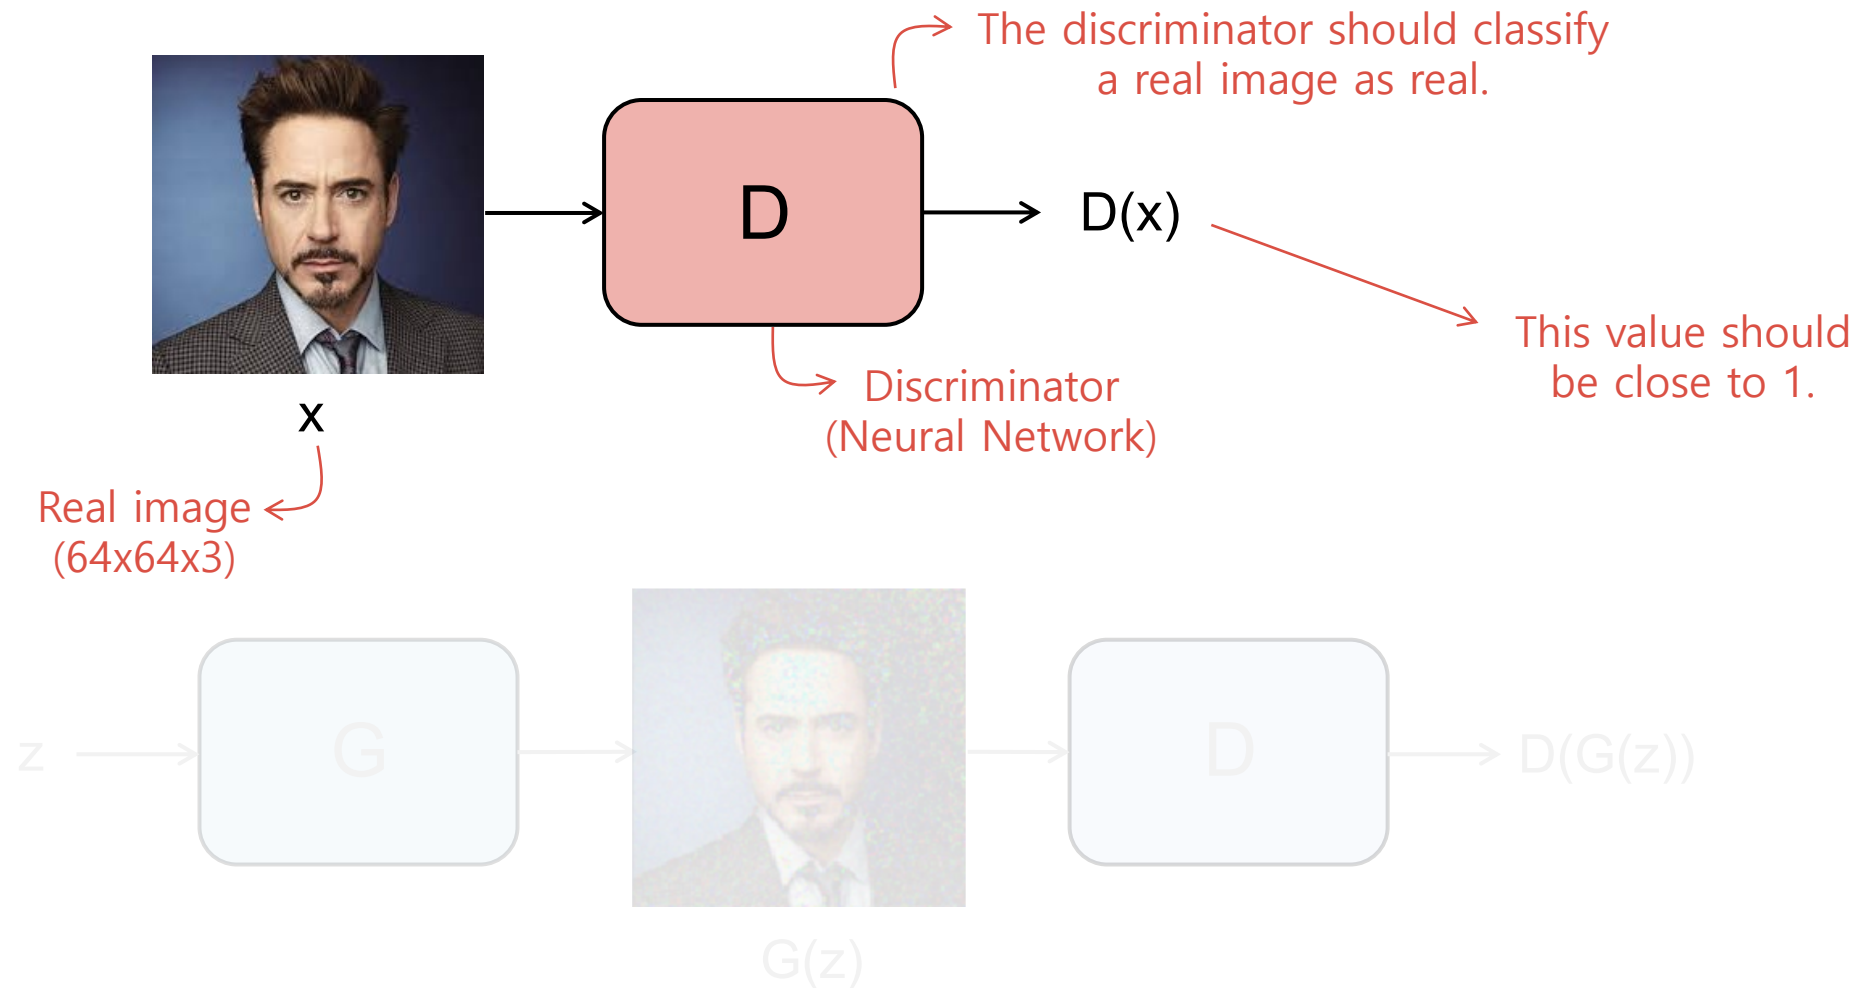
GANs



Intuition in GAN



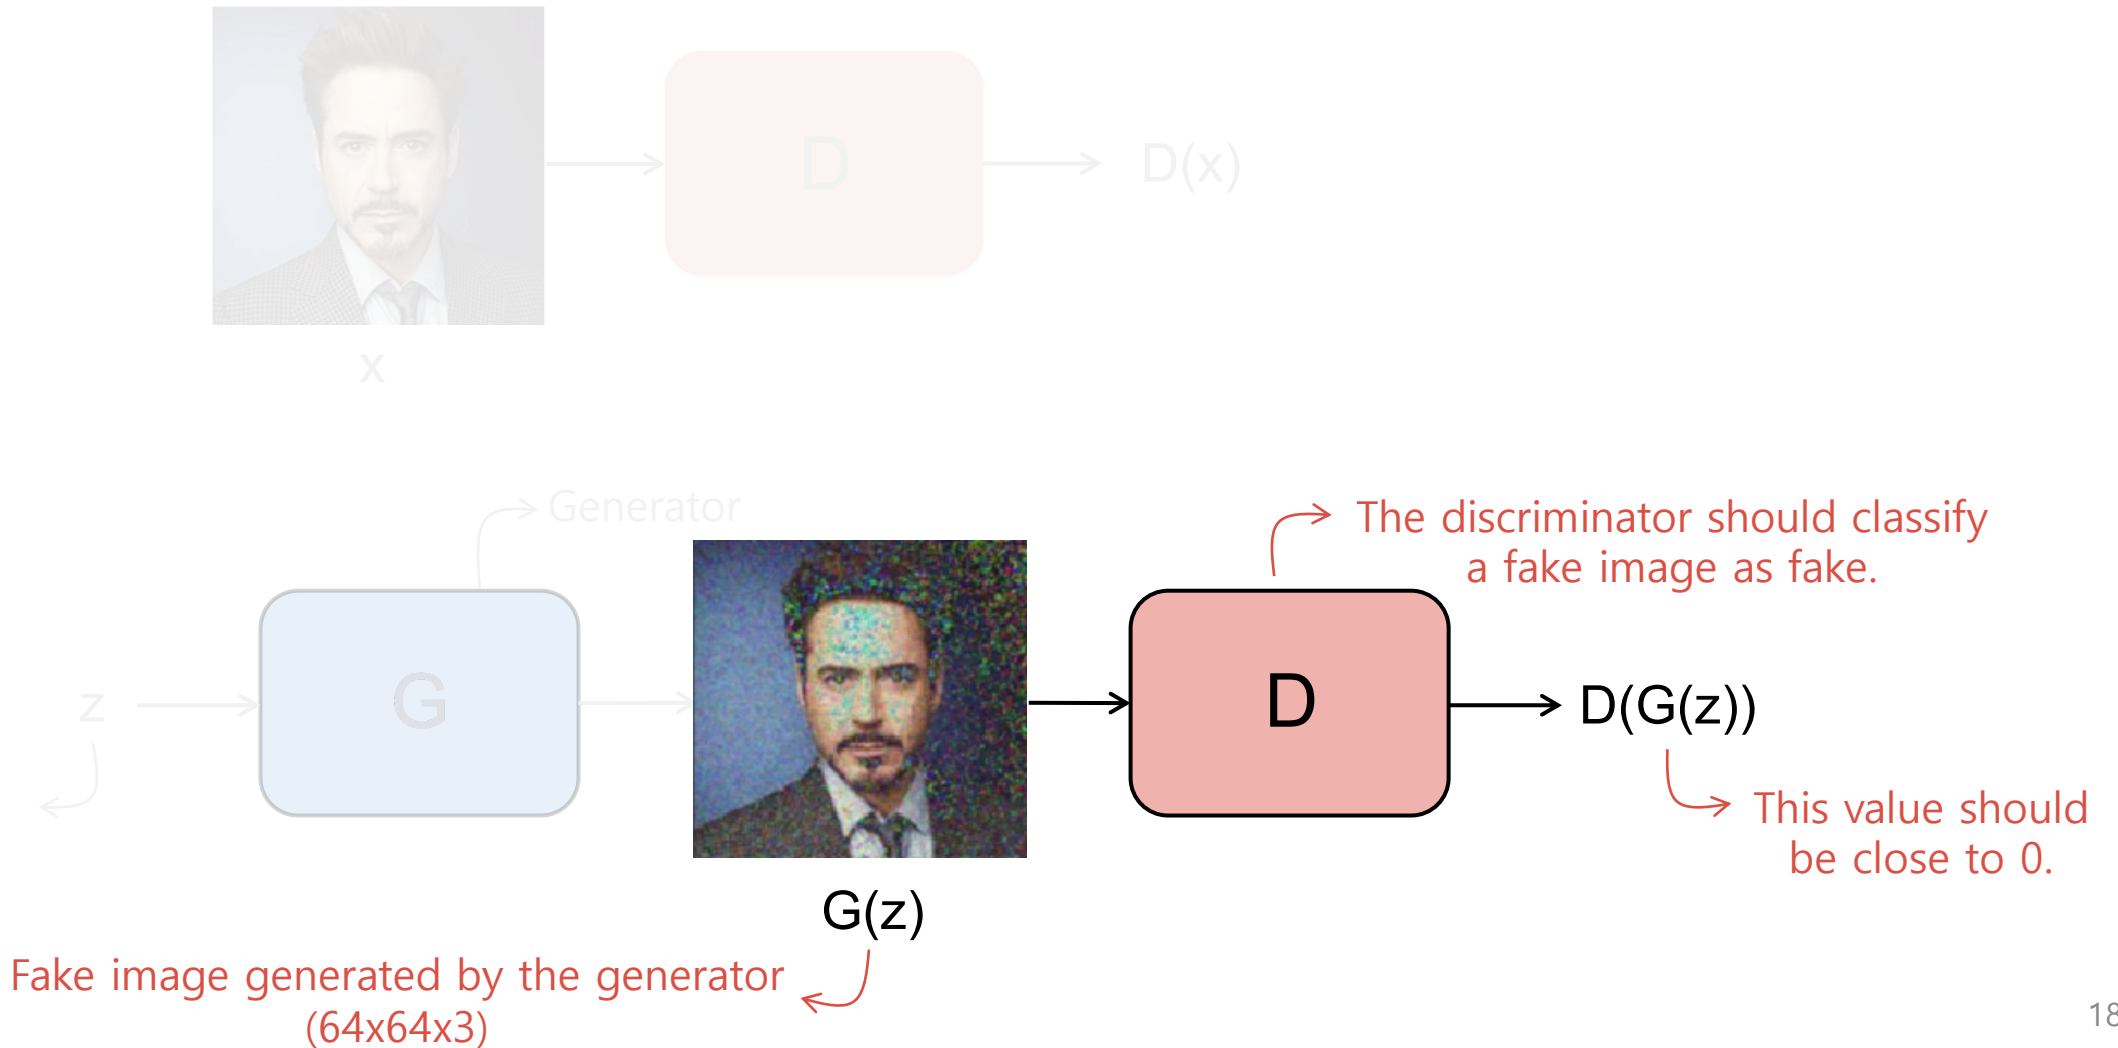
GANs



Intuition in GAN



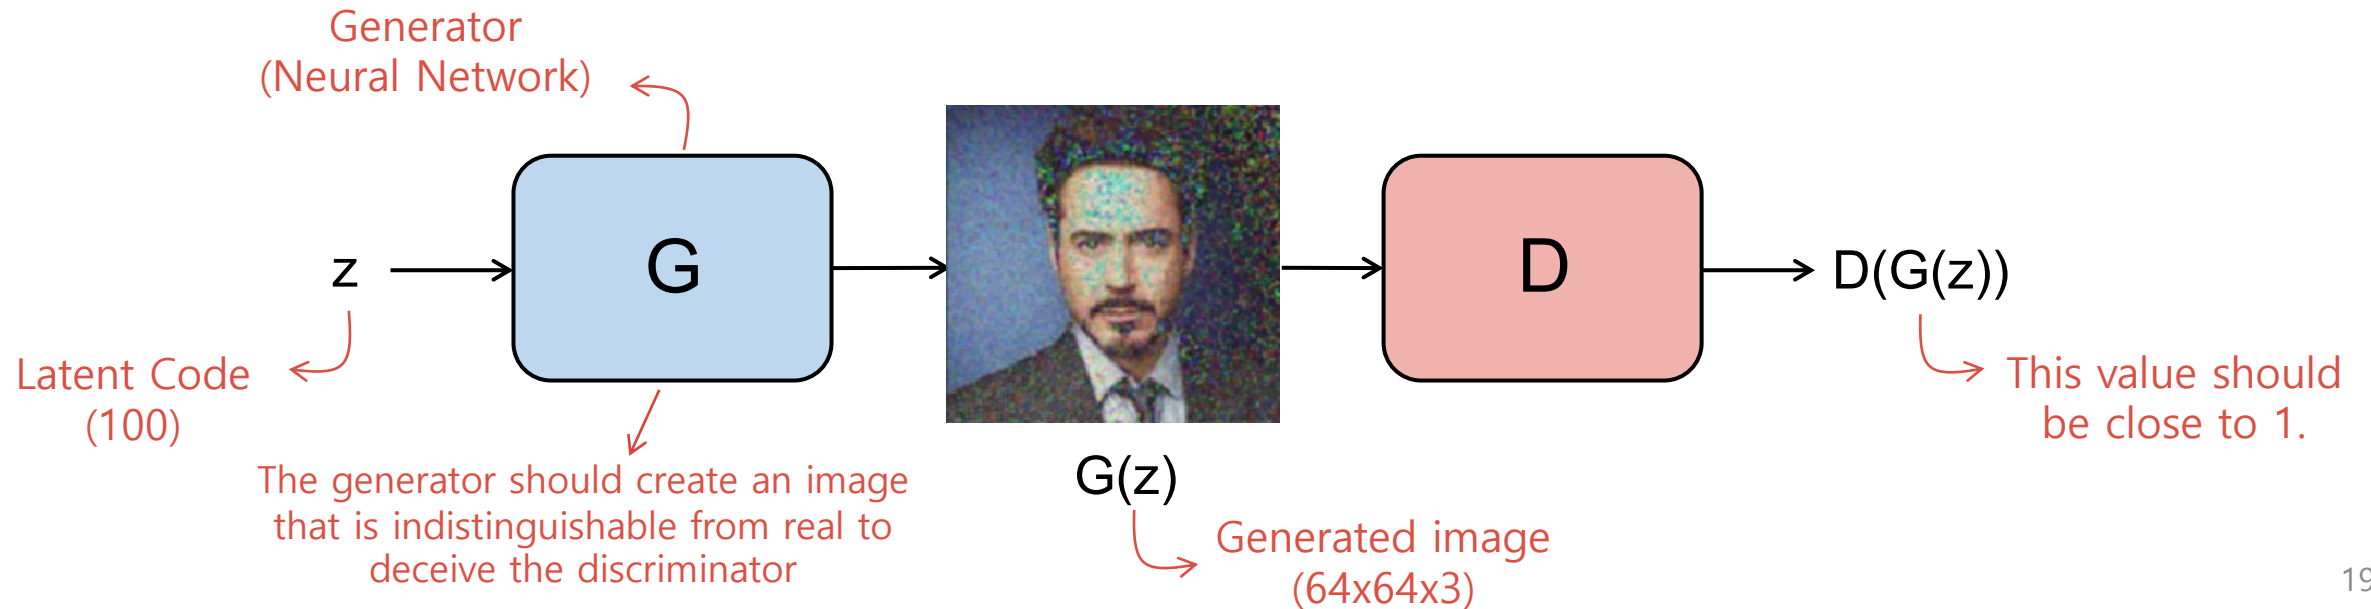

GANs



Intuition in GAN



GANs



Objective Function of GAN



GANs



Sample x from real data distribution

Sample latent code z from Gaussian distribution

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)} [\log D(x)] + E_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

D should maximize $V(D, G)$


Maximum when $D(x) = 1$

Maximum when $D(G(z)) = 0$

Objective function




x



$D(x)$

Train D to classify real images as real

Training with real images



Objective Function of GAN



GANs



$$\min_G \max_D V(D, G) = \cancel{E_{x \sim p_{data}(x)} [\log D(x)]} + E_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

(Note: The first term is crossed out in the original image. An arrow points to it with the text "G is independent of this part".)


G should minimize $V(D, G)$

Minimum when $D(G(z)) = 1$

Objective function




X




D(x)


Training with real images



Train G to deceive D




G(z)




D(G(z))

Training with fake images



Training with real images



Training with fake images



```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```




Define the discriminator

input size: 784
hidden size: 128
output size: 1

Assume x is MNIST (784 dimension)




```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```




Define the generator


input size: 100
hidden size: 128
output size: 784





```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```

Binary Cross Entropy Loss $(h(x), y)$
 $-y \log h(x) - (1 - y) \log(1 - h(x))$





```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```



Optimizer for D and G

```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```




x is a tensor of shape (batch_size, 784).
 z is a tensor of shape (batch_size, 100).


```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```



Train the discriminator
with real images




Train the discriminator
with fake images




Forward, Backward and Gradient Descent

```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```





Train the generator
to deceive the discriminator



Forward, Backward and Gradient Descent →

```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```



The complete code can be found here


<https://github.com/yunjey/pytorch-tutorial>

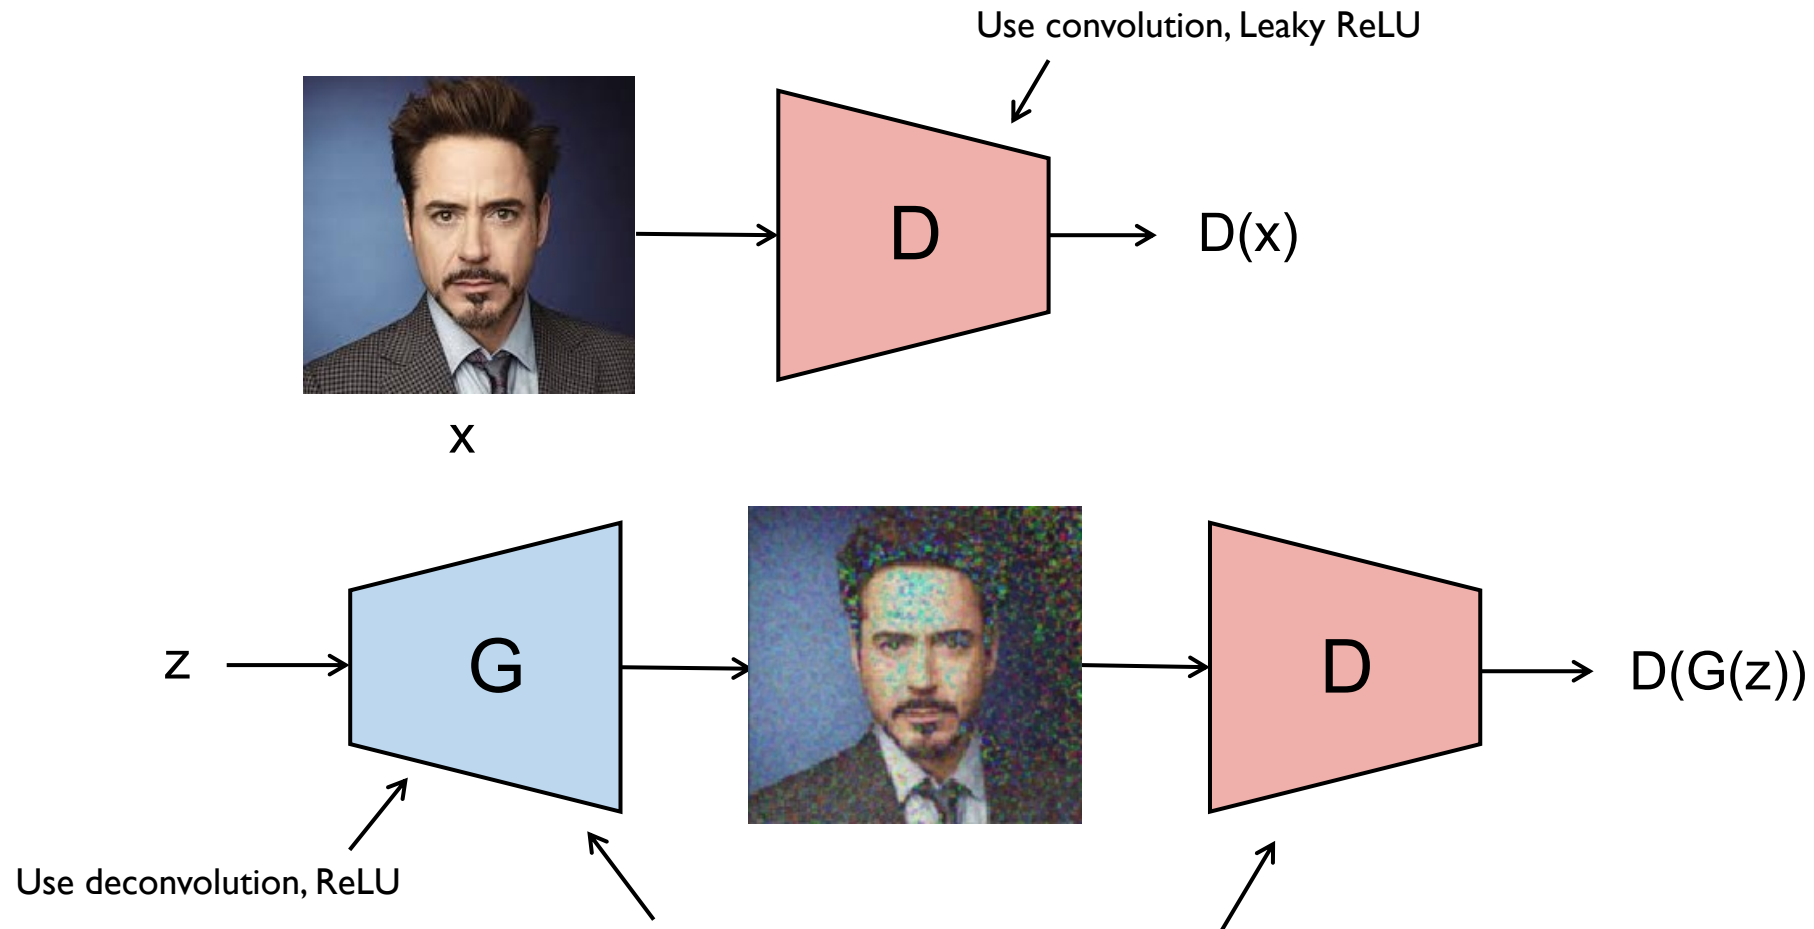
```
1 import torch
2 import torch.nn as nn
3
4
5 D = nn.Sequential(
6     nn.Linear(784, 128),
7     nn.ReLU(),
8     nn.Linear(128, 1),
9     nn.Sigmoid())
10
11 G = nn.Sequential(
12     nn.Linear(100, 128),
13     nn.ReLU(),
14     nn.Linear(128, 784),
15     nn.Tanh())
16
17 criterion = nn.BCELoss()
18
19 d_optimizer = torch.optim.Adam(D.parameters(), lr=0.01)
20 g_optimizer = torch.optim.Adam(G.parameters(), lr=0.01)
21
22 # Assume x be real images of shape (batch_size, 784)
23 # Assume z be random noise of shape (batch_size, 100)
24
25 while True:
26     # train D
27     loss = criterion(D(x), 1) + criterion(D(G(z)), 0)
28     loss.backward()
29     d_optimizer.step()
30
31     # train G
32     loss = criterion(D(G(z)), 1)
33     loss.backward()
34     g_optimizer.step()
```



- Deep Convolutional GAN(DCGAN), 2015

The authors present a model that is still highly preferred.






- No pooling layer (Instead strided convolution)
- Use batch normalization
- Adam optimizer($\text{lr}=0.0002, \text{beta1}=0.5, \text{beta2}=0.999$)



- Latent vector arithmetic



04

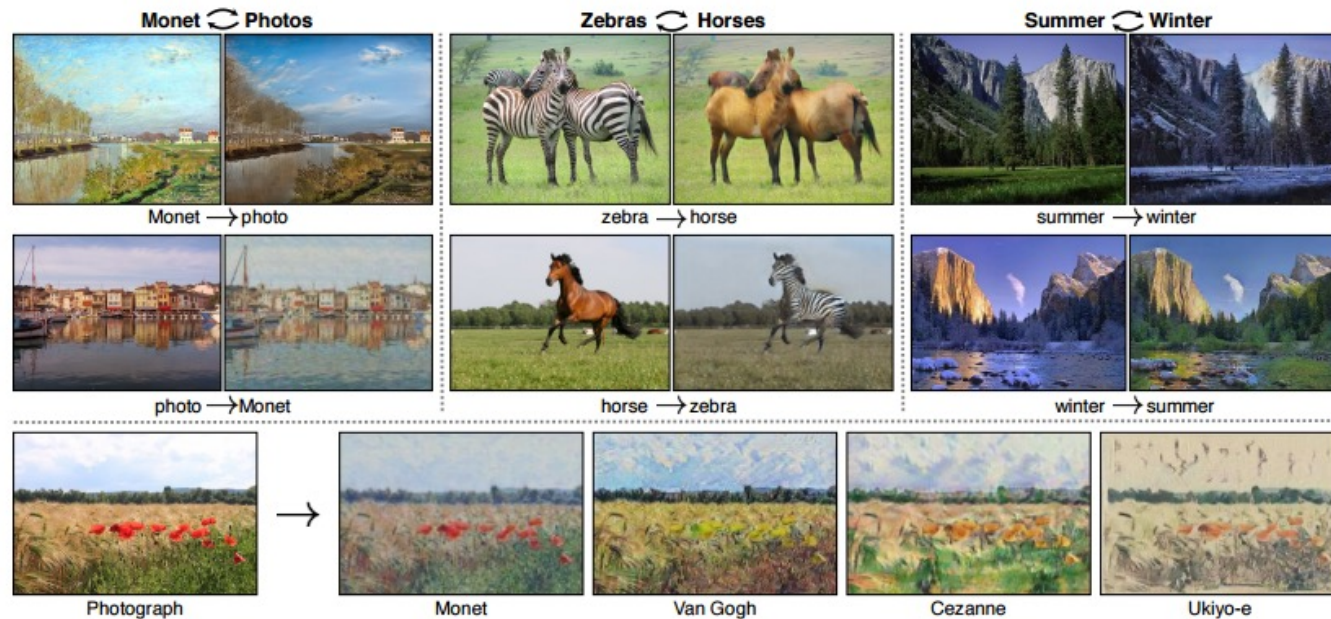
Extensions of GAN





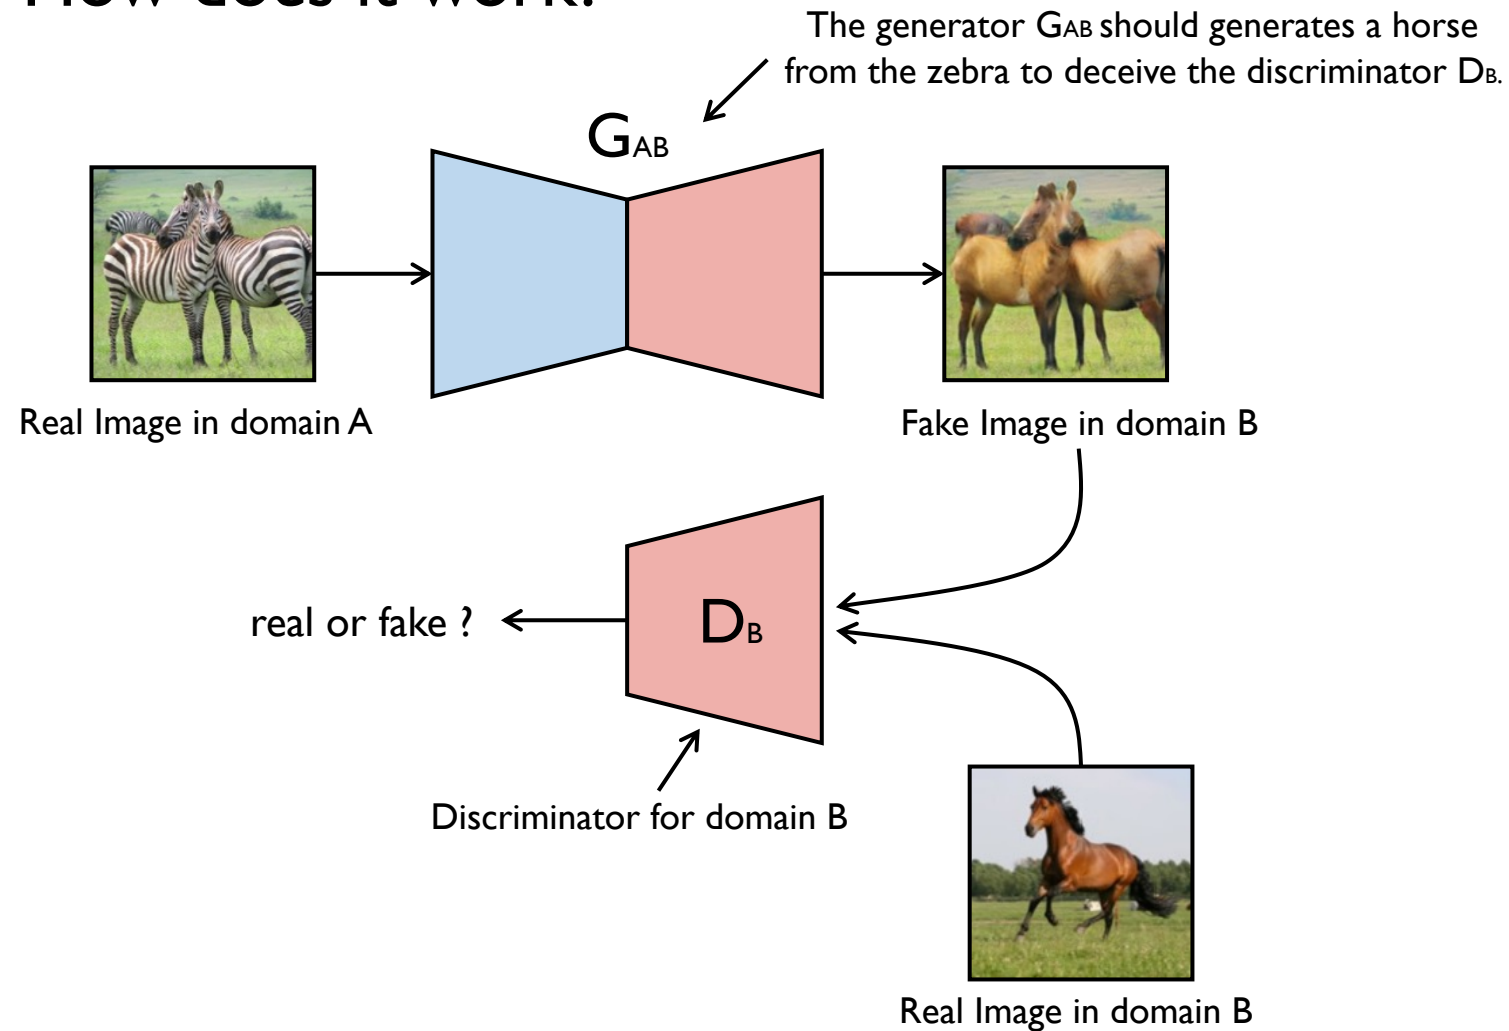
- CycleGAN: Unpaired Image-to-Image Translation

presents a GAN model that transfer an image from a source domain A to a target domain B in the absence of paired examples.






- How does it work?



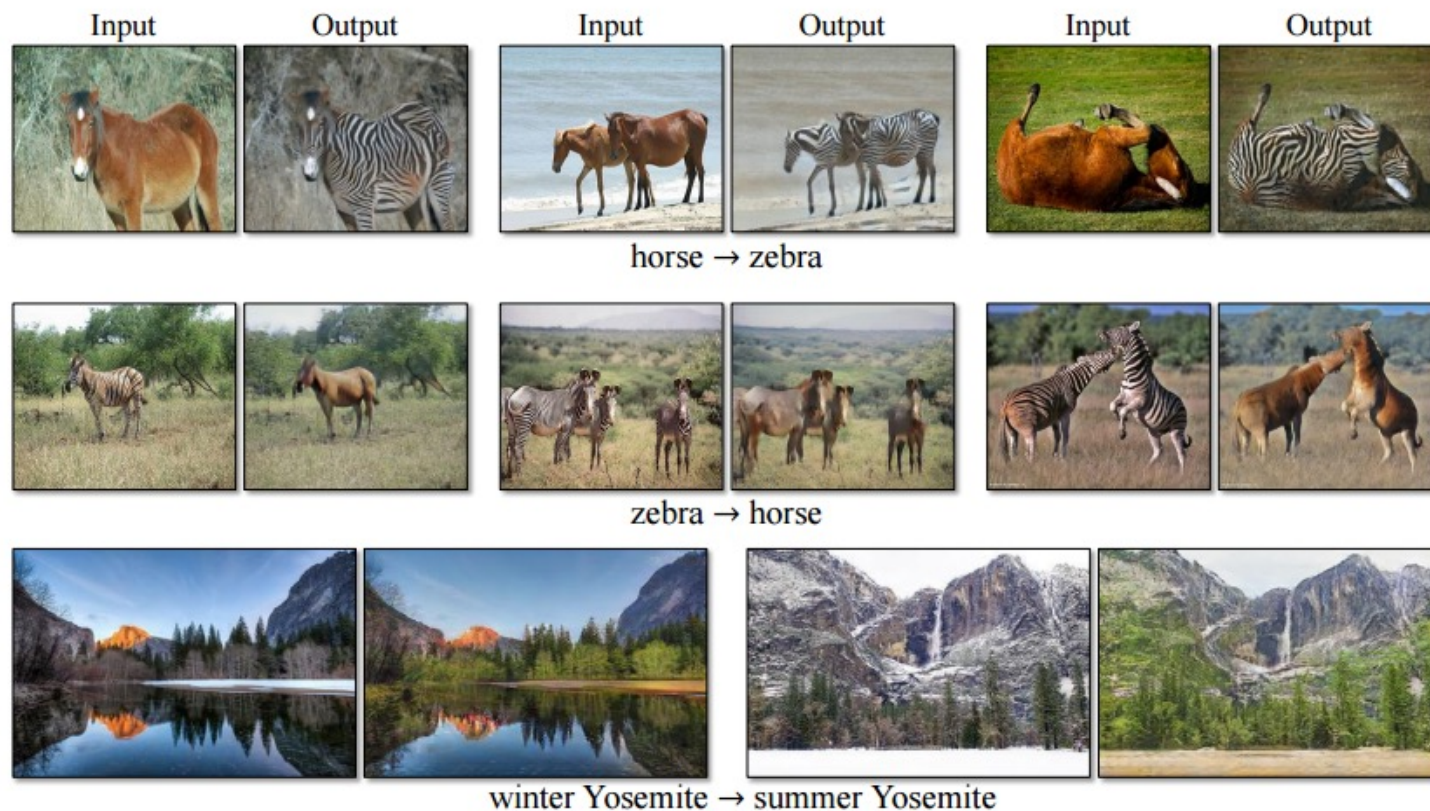


- How does it work?





- Results






- Results

Odd columns contain real images and even columns contain generated images.



SVHN-to-MNIST

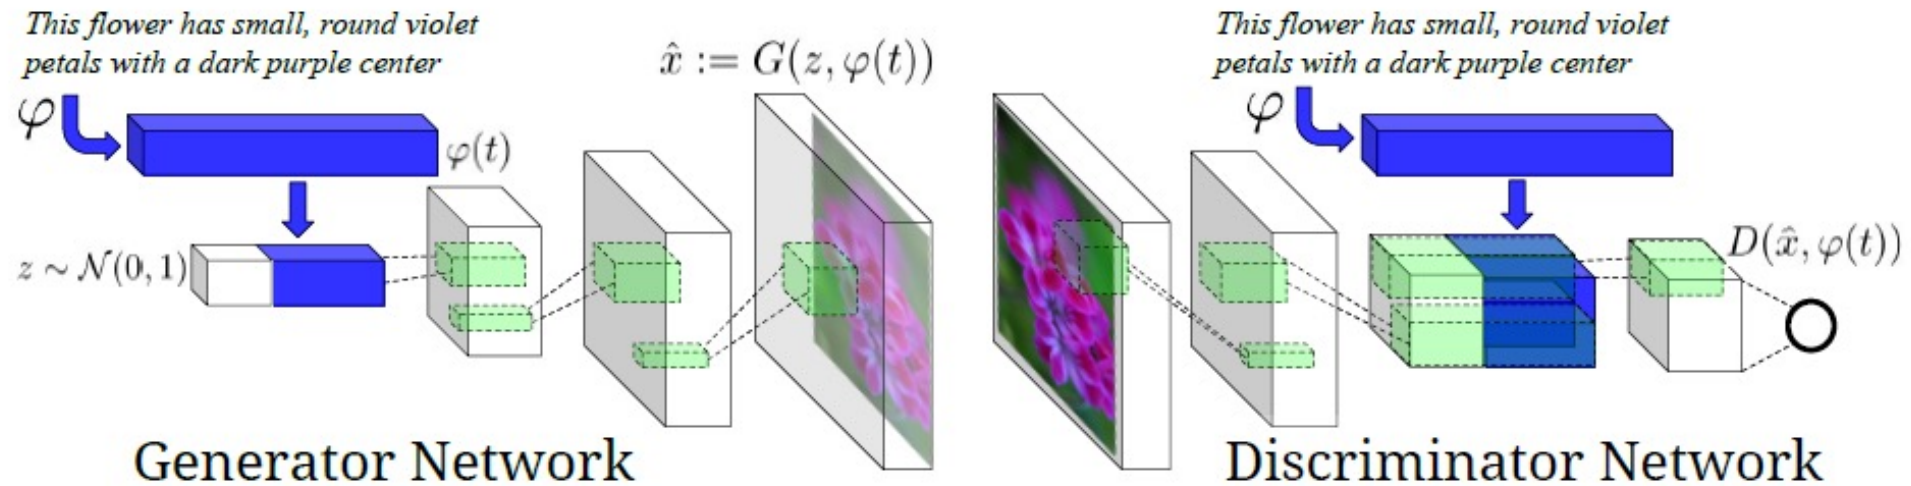


MNIST-to-SVHN



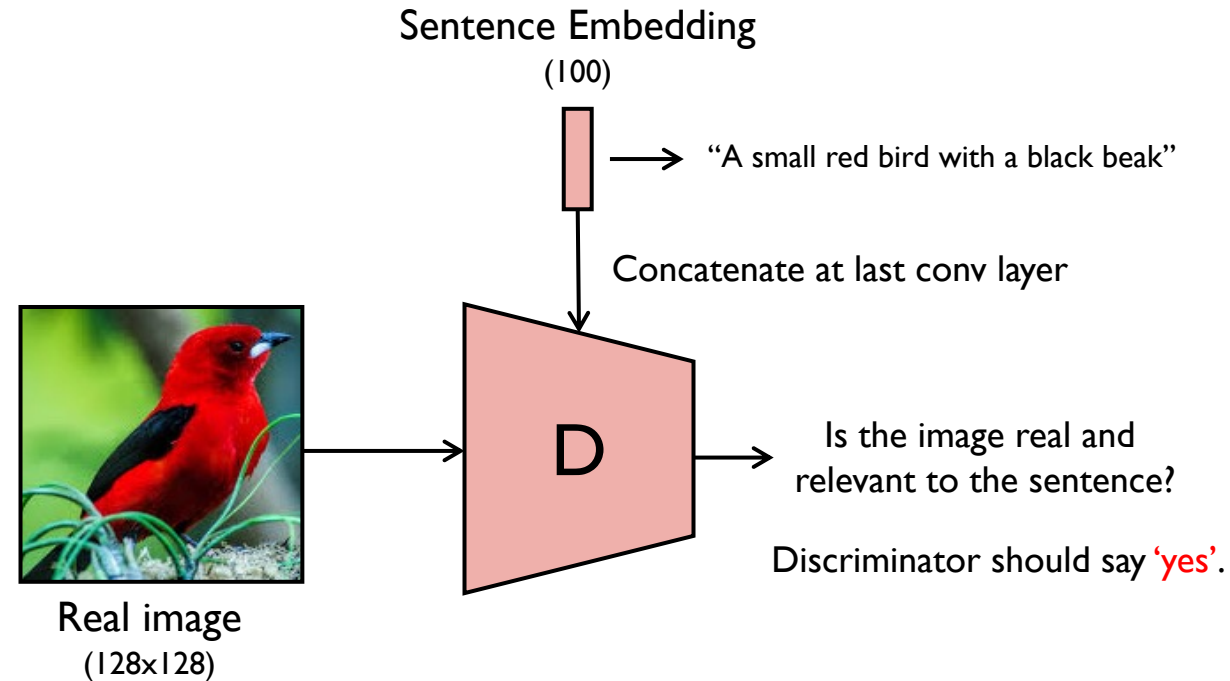
- Generative Adversarial Text to Image Synthesis, 2016

presents a novel model architecture that generates an image from the text.



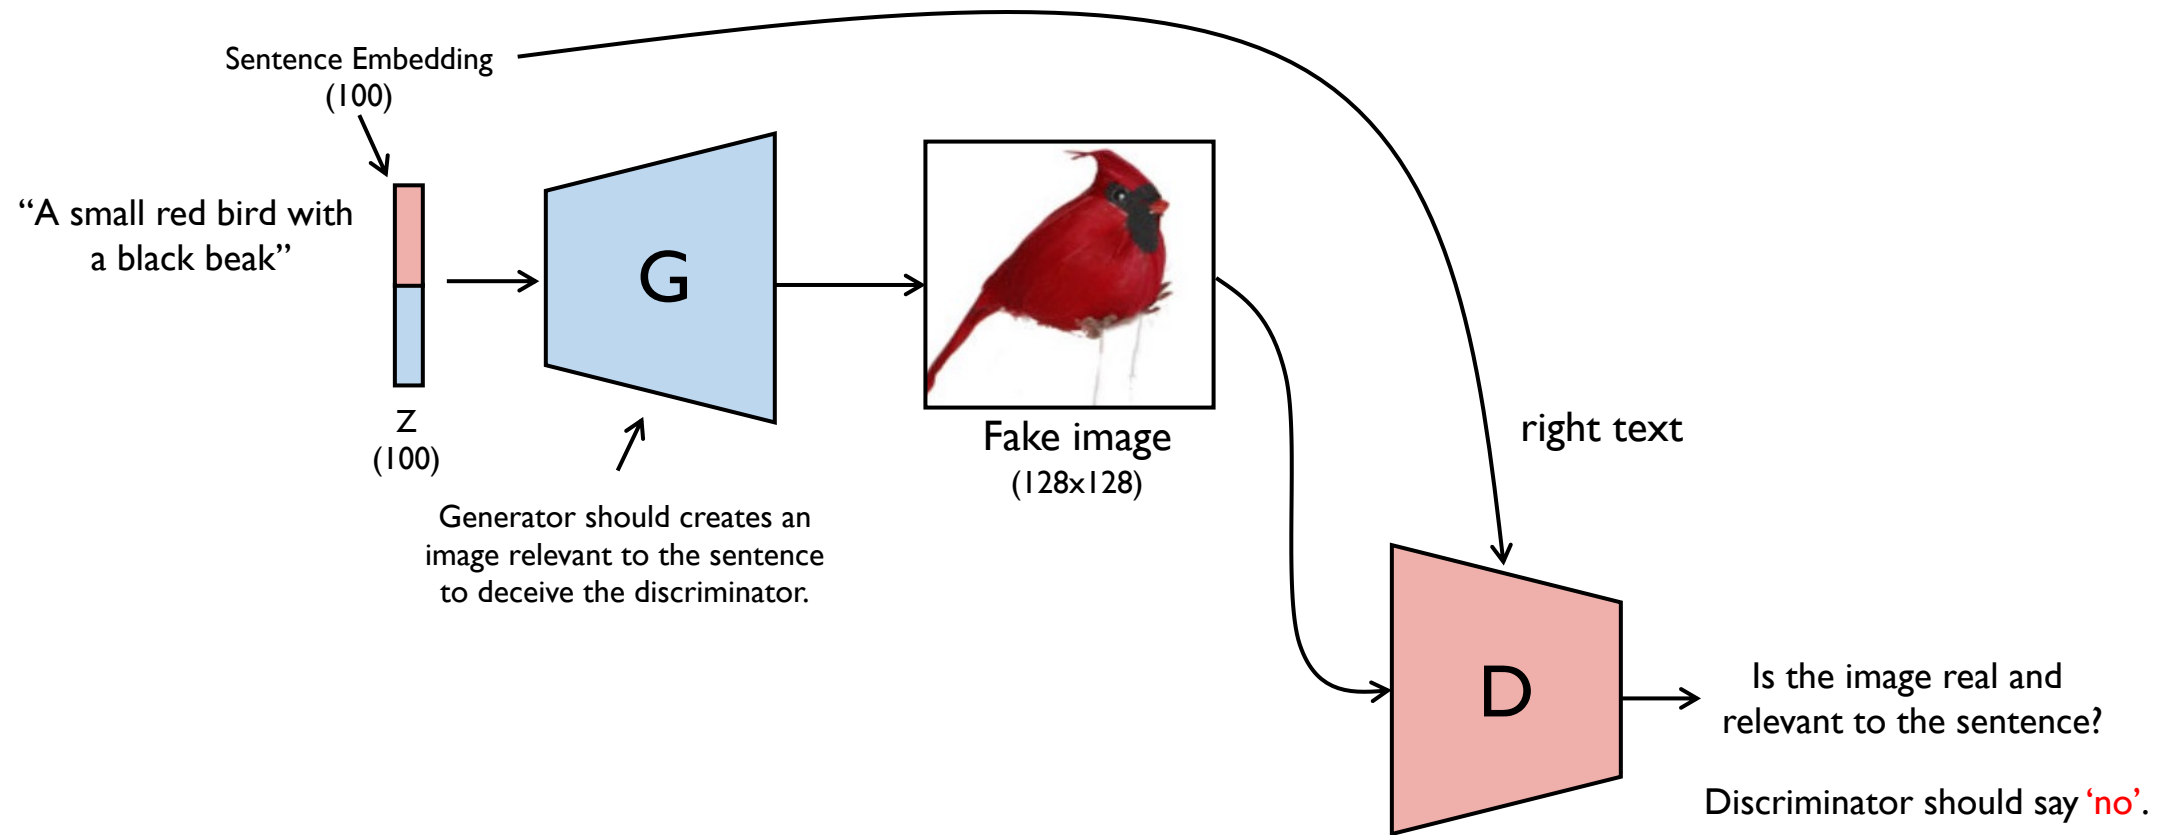


- Training with (real image, right text)



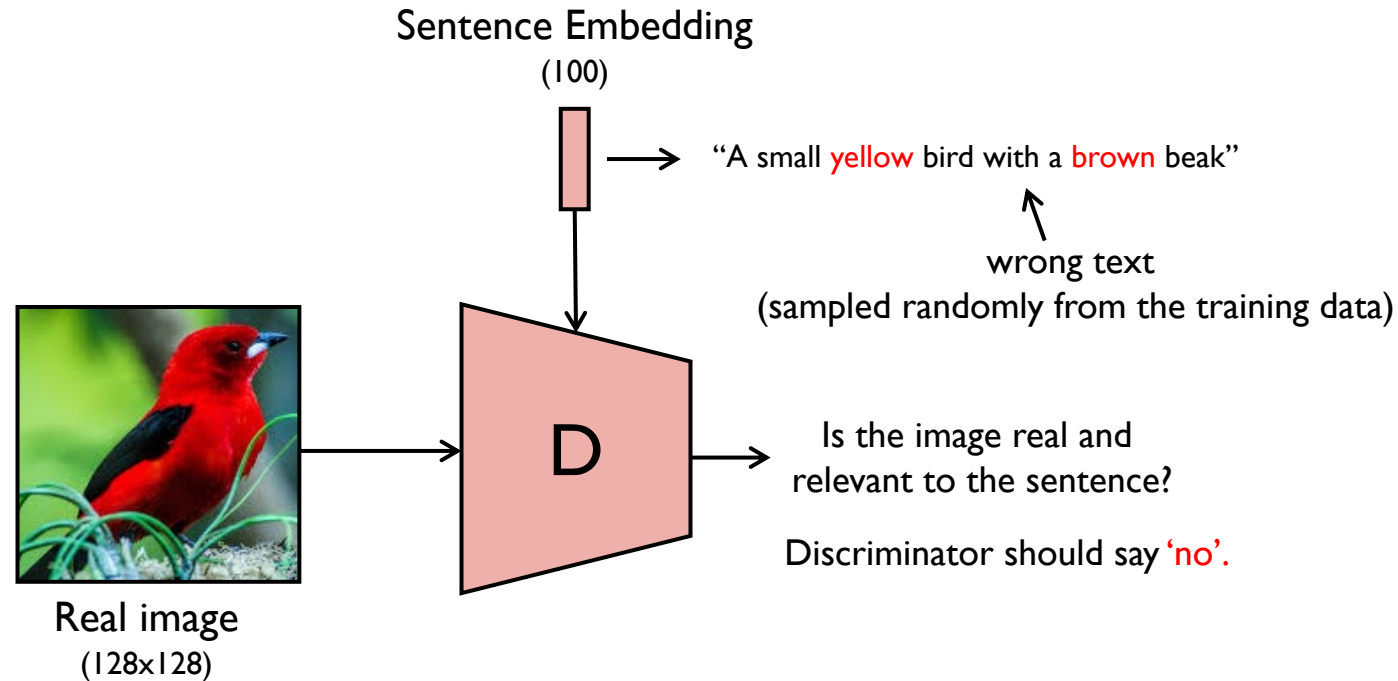


- Training with (**fake image**, right text)



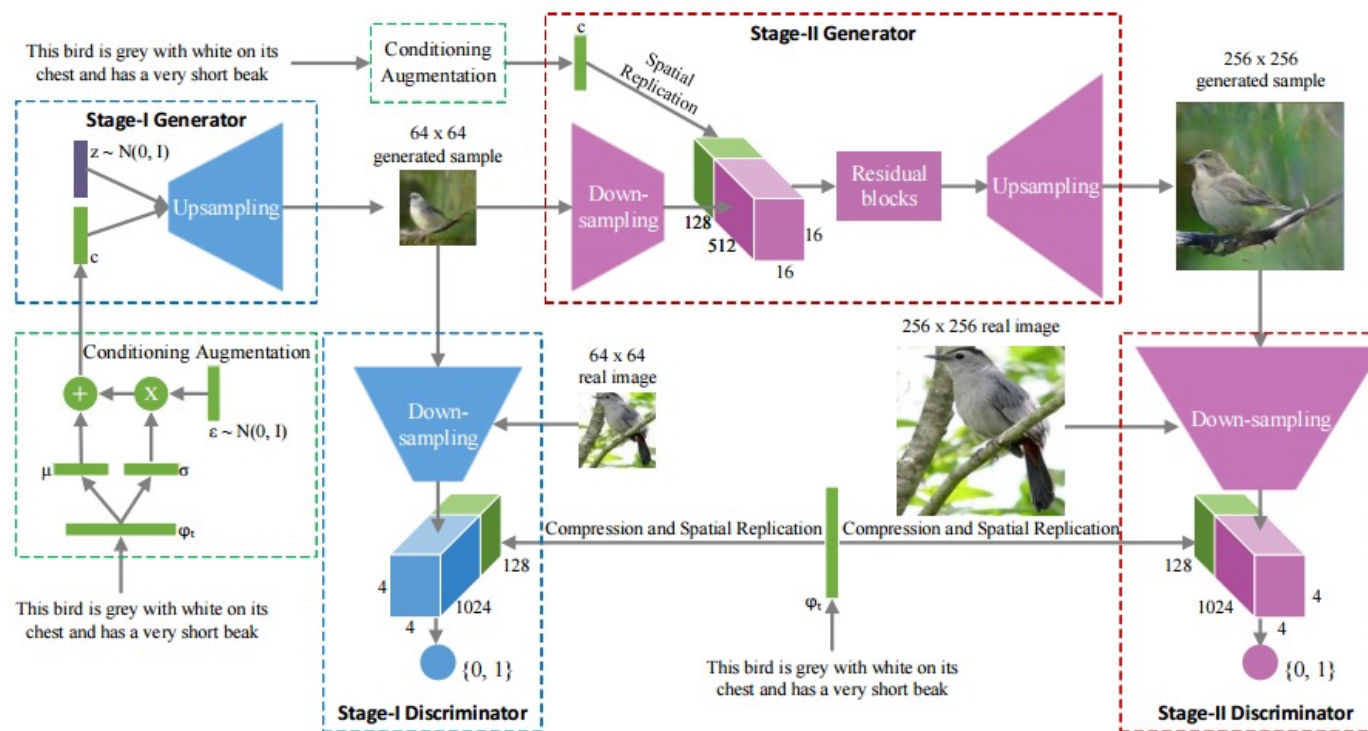


- Training with (real image, **wrong text**)





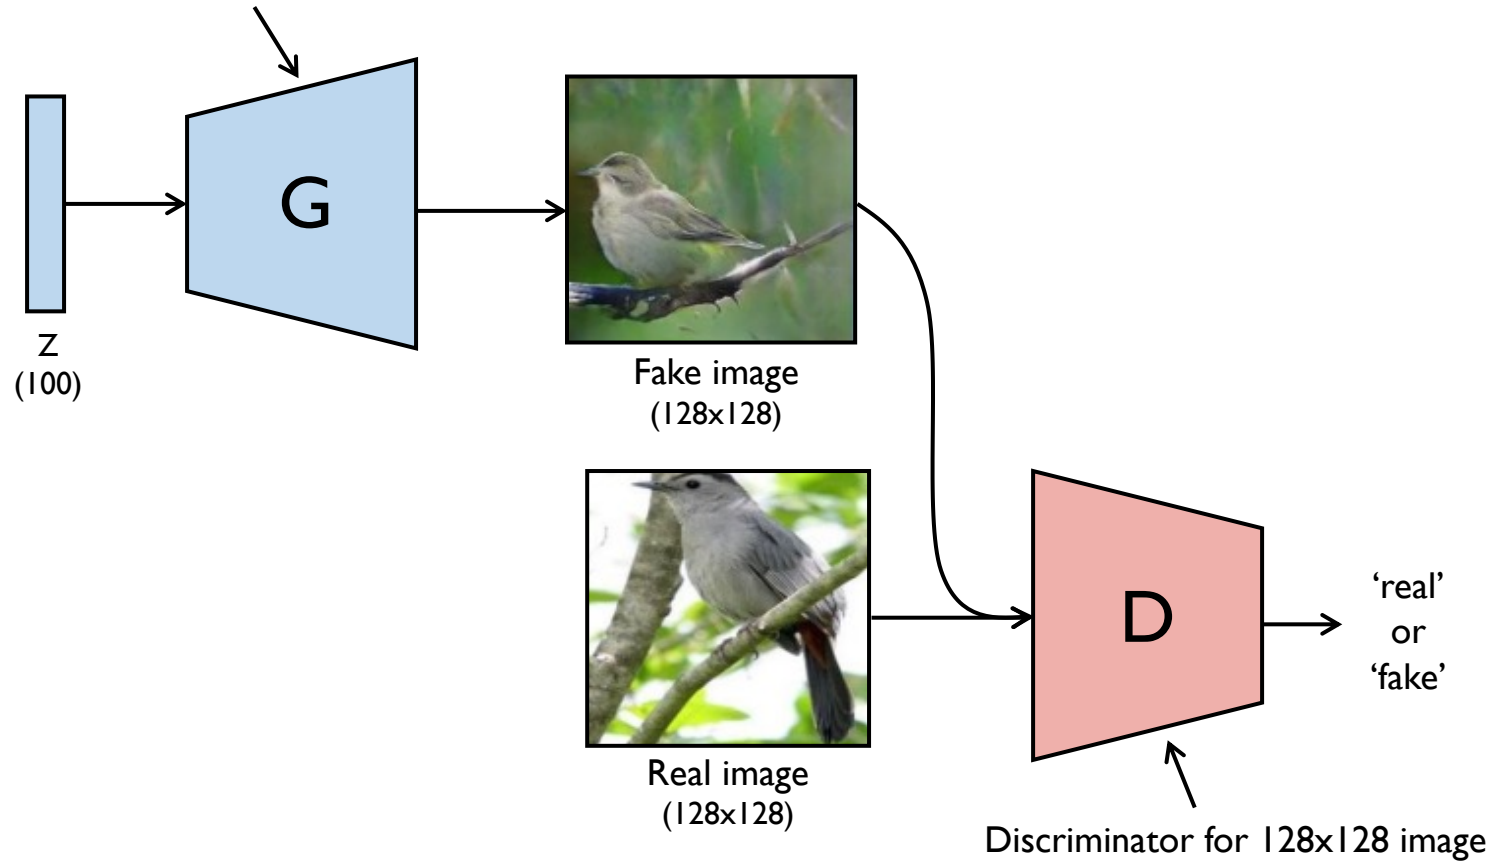
- StackGAN: Text to Photo-realistic Image Synthesis





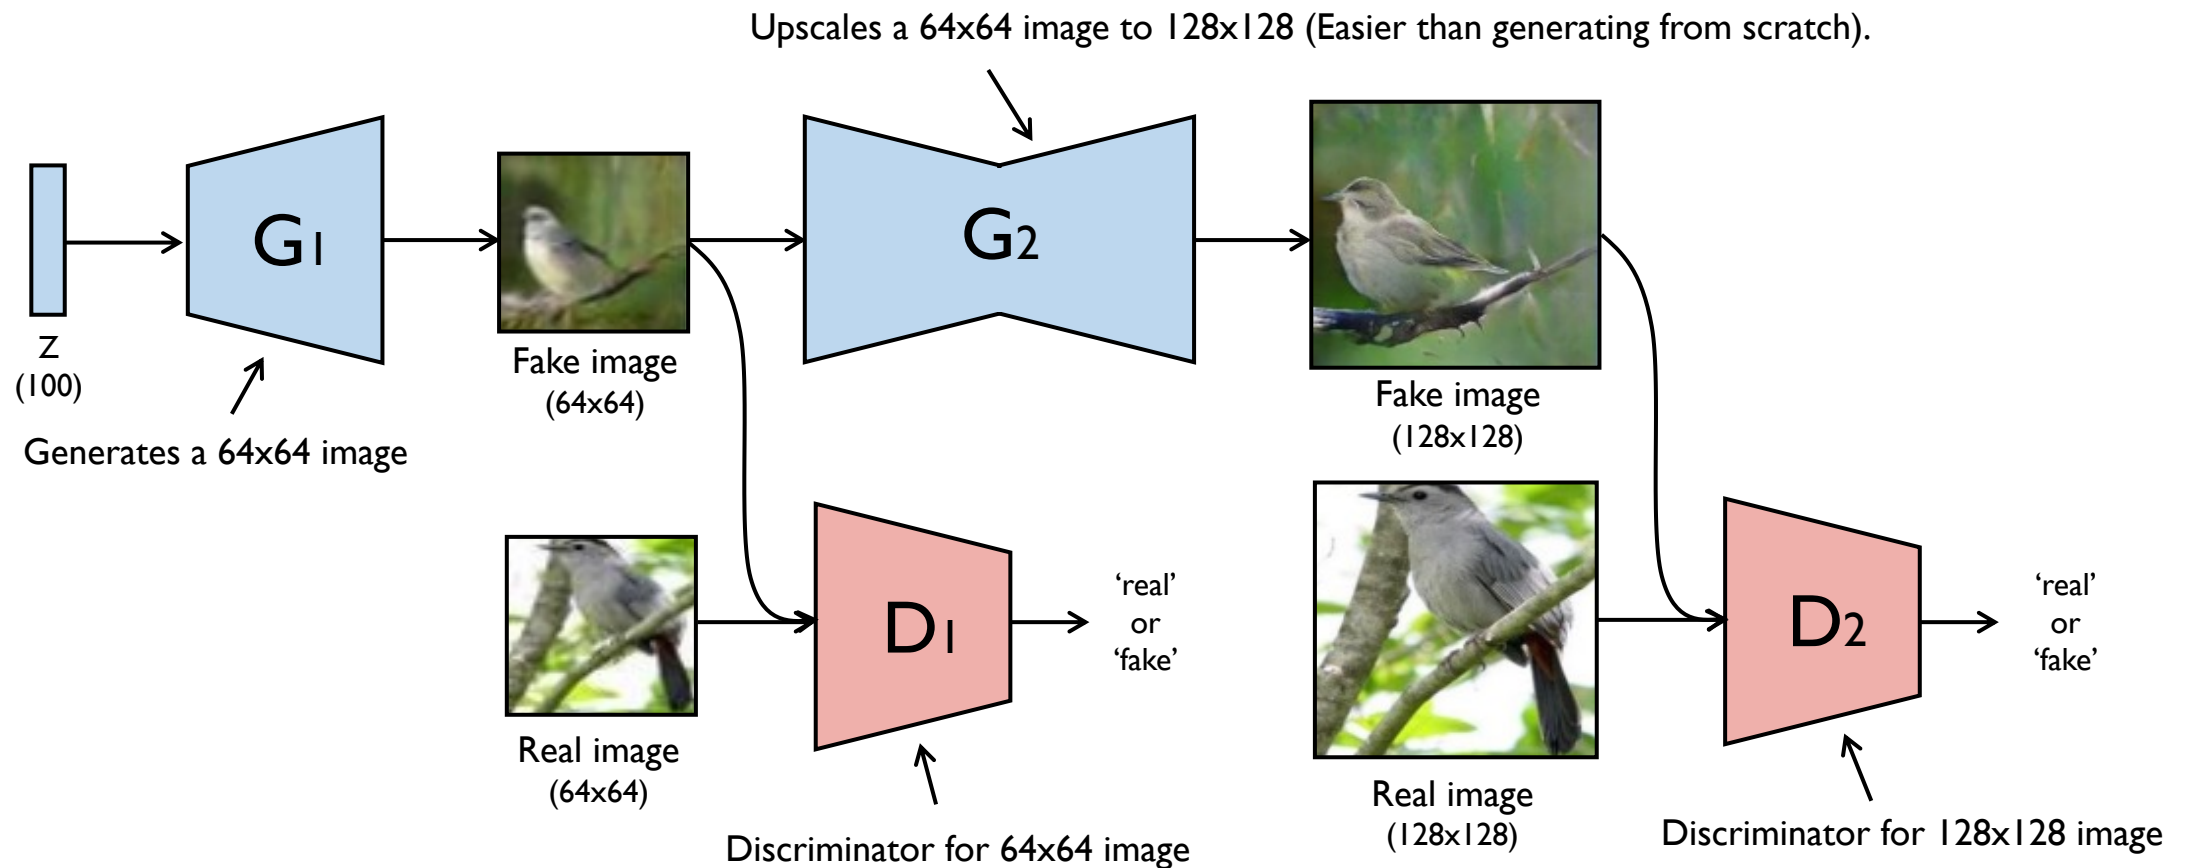
- Generating 128x128 from scratch

Generates a 128x128 image from scratch (not guarantee good result)





- Generating 128x128 from 64x64



Thank you

